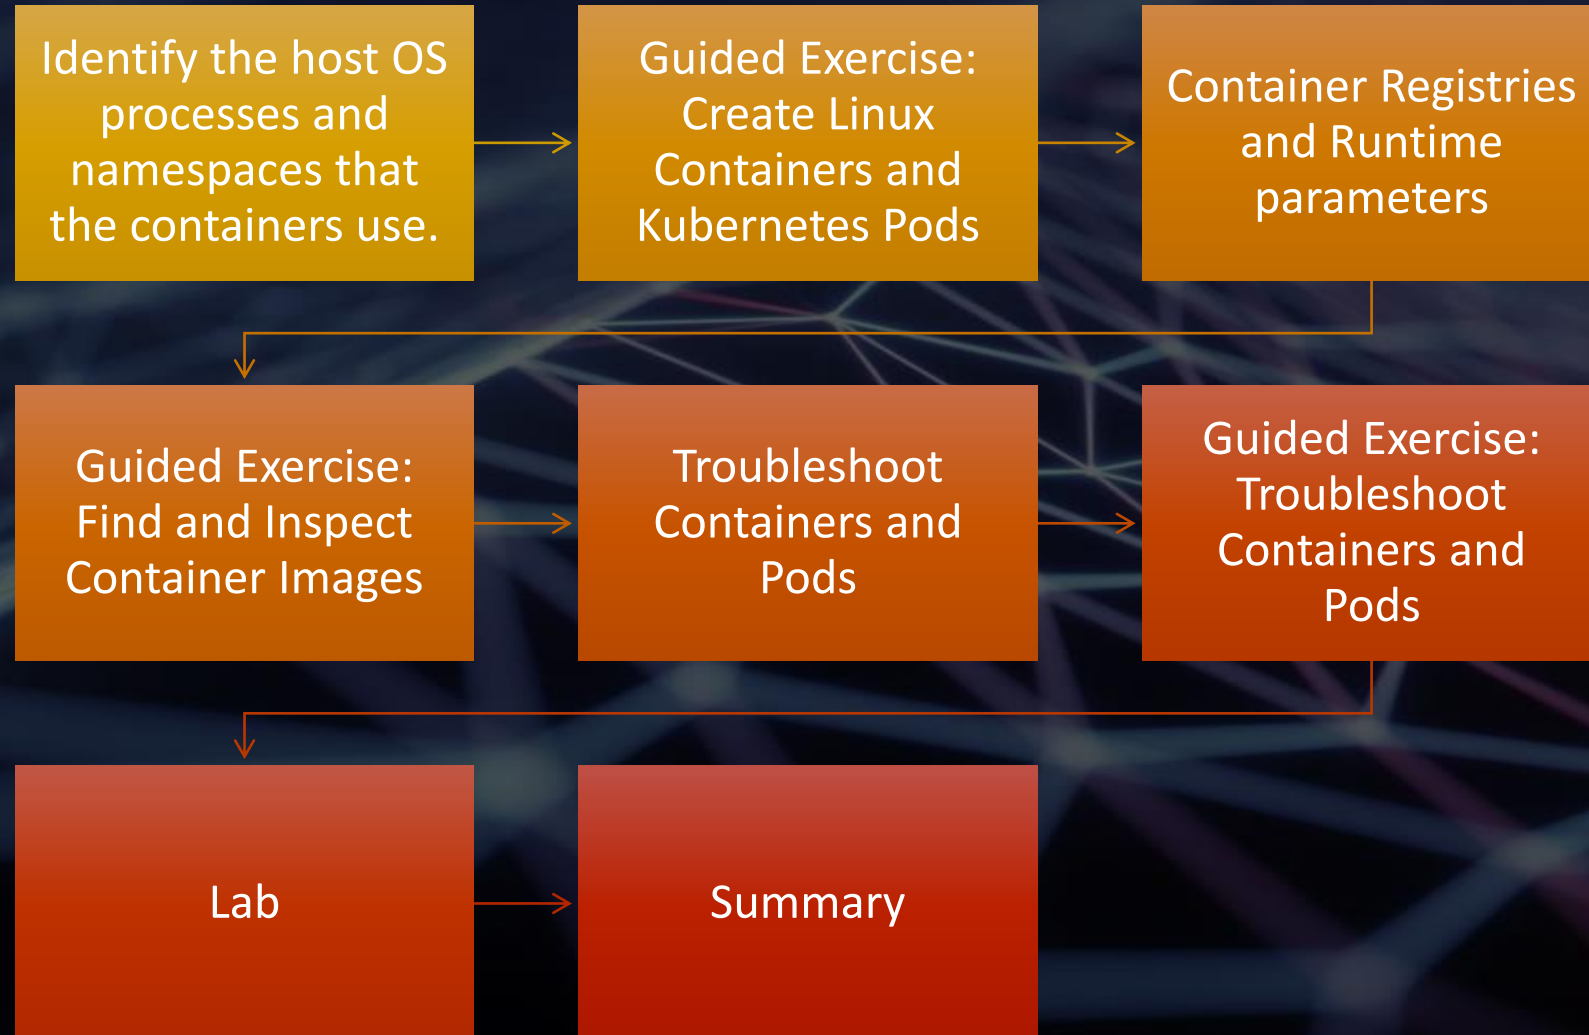


Chapter 3: Run Applications as Containers and Pods



Chapter objectives



Lesson 1 : Create Linux Containers and Kubernetes Pods

- Run containers inside pods and identify the host OS processes and namespaces that the containers use

Creating Containers and Pods

- Syntax

```
oc run RESOURCE/NAME --image IMAGE [options] or  
kubectl run RESOURCE/NAME --image IMAGE [options]
```

- Example

```
oc run web-server --image registry.access.redhat.com/ubi8/httpd-24  
kubectl run my-app --image registry.access.redhat.com/ubi9/ubi
```

Runtime Parameters / Options

- --command
- -i / --stdin
- -t / --tty
- --restart
- --env
- --rm

Use the --command option

- Syntax

```
oc run RESOURCE/NAME --image IMAGE --command -- cmd arg1 ... argN
```

```
kubectl run RESOURCE/NAME --image IMAGE --command -- cmd arg1 ... argN
```

- Example

```
oc run web-server --image registry.access.redhat.com/ubi8/httpd-24 --command -- sleep 500
```

```
kubectl run my-app --image registry.access.redhat.com/ubi9/ubi --command -- date
```

Use the --rm and -i -t option

- -i option : Interactive mode --rm option : delete pod after it exits
- -t option : open a TTY session --restart : restart pod
- Example 1

```
[user@host ~]$ oc run -it my-app \
--image registry.access.redhat.com/ubi9/ubi \
--restart Never --command -- date
Mon Feb 20 22:36:55 UTC 2023
```

- Example 2

```
[user@host ~]$ kubectl run -it my-app --rm \
--image registry.access.redhat.com/ubi9/ubi \
--restart Never --command -- date
Mon Feb 20 22:38:50 UTC 2023
pod "date" deleted
```


The --restart option

Accepted Values	Description
--restart Always	Cluster continuously tries to restart a successfully exited container, for up to five minutes. The default pod restart policy is Always.
--restart OnFailure	Setting the pod restart policy to OnFailure tells the cluster to restart only failed containers in the pod, for up to five minutes
--restart Never	If the restart policy is set to Never, then the cluster does not try to restart exited or failed containers in a pod. Instead, the pods immediately fail and exit.

The --env option

- Start a database system using MySQL image. At same time set the root password

```
[user@host ~]$ oc run mysql \
--image registry.redhat.io/rhel9/mysql-80 \
--env MYSQL_ROOT_PASSWORD=myP$$123
pod/mysql created
```

User and Group IDs Assignment

- Defined at project/namespace level
- Determine UID and GID for pods and containers
- Default security policies
 - rootless container cannot choose USER or UID
 - OpenShift assigns user in container to UID and GID from range
 - ignore USER instruction in container image
 - rootfull container can choose USER - but not recommended
- Red Hat recommendations:
 - Always run containers as rootless
 - Use unprivileged user with necessary privileges to specific commands

The UID and supplemental GID range

```
[user@host ~]$ oc describe project my-app
Name:      my-app
...output omitted...
Annotations:  openshift.io/description=
              openshift.io/display-name=
              openshift.io/requester=developer
              openshift.io/sa.scc.mcs=s0:c27,c4
              openshift.io/sa.scc.supplemental-groups=1000710000/10000
              openshift.io/sa.scc.uid-range=1000710000/10000
...output omitted...
```

- The UID and GID range follow the format
<first_id>/<id_pool_size> or <first_id>-<last_id>
- Assign distinct range of UIDs and GIDs for each project
 - ensures applications do not run as the same UID and GID
 - if two containers is the same, processes in one container can access resources in the other

Pod Security

- Security Context Constraints (SCC)
- *Grant* or *Deny* OS-level privileges in pods
- Undefined security context - issue warning
- Safely ignore pod security warnings in course exercises
- Discussed in more details in DO280 course.

Execute Commands in Running Containers

- Use the **oc exec** or **kubectl exec** command
- The syntax

```
oc exec RESOURCE/NAME -- COMMAND [args...] [options]
```

- Example: Execute date command on running my-app

```
[user@host ~]$ oc exec my-app -- date  
Tue Feb 21 20:43:53 UTC 2023
```

- Example: For multiple container in pod, use the -c option

```
[user@host ~]$ kubectl exec my-app -c ruby-container -- date  
Tue Feb 21 20:46:50 UTC 2023
```

Execute Commands in Running Containers

- Combine with -i and -t options

```
[user@host ~]$ oc exec my-app -c ruby-container -it -- bash -il  
[1000780000@ruby-container /]$
```

- Exit command to terminate the interactive session

```
[user@host ~]$ kubectl exec my-app -c ruby-container -it -- bash -il  
[1000780000@ruby-container /]$ date  
Tue Feb 21 21:16:00 UTC 2023  
[1000780000@ruby-container] exit
```

Attach to a running container

- is to gain access to the input/output stream of a container.

```
# Get output from running pod mypod; use the 'oc.kubernetes.io/default-container' annotation  
# for selecting the container to be attached or the first container in the pod will be chosen  
oc attach mypod
```

```
# Get output from ruby-container from pod mypod  
oc attach mypod -c ruby-container
```

```
# Switch to raw terminal mode; sends stdin to 'bash' in ruby-container from pod mypod  
# and sends stdout/stderr from 'bash' back to the client  
oc attach mypod -c ruby-container -i -t
```

```
# Get output from the first pod of a replica set named nginx  
oc attach rs/nginx
```

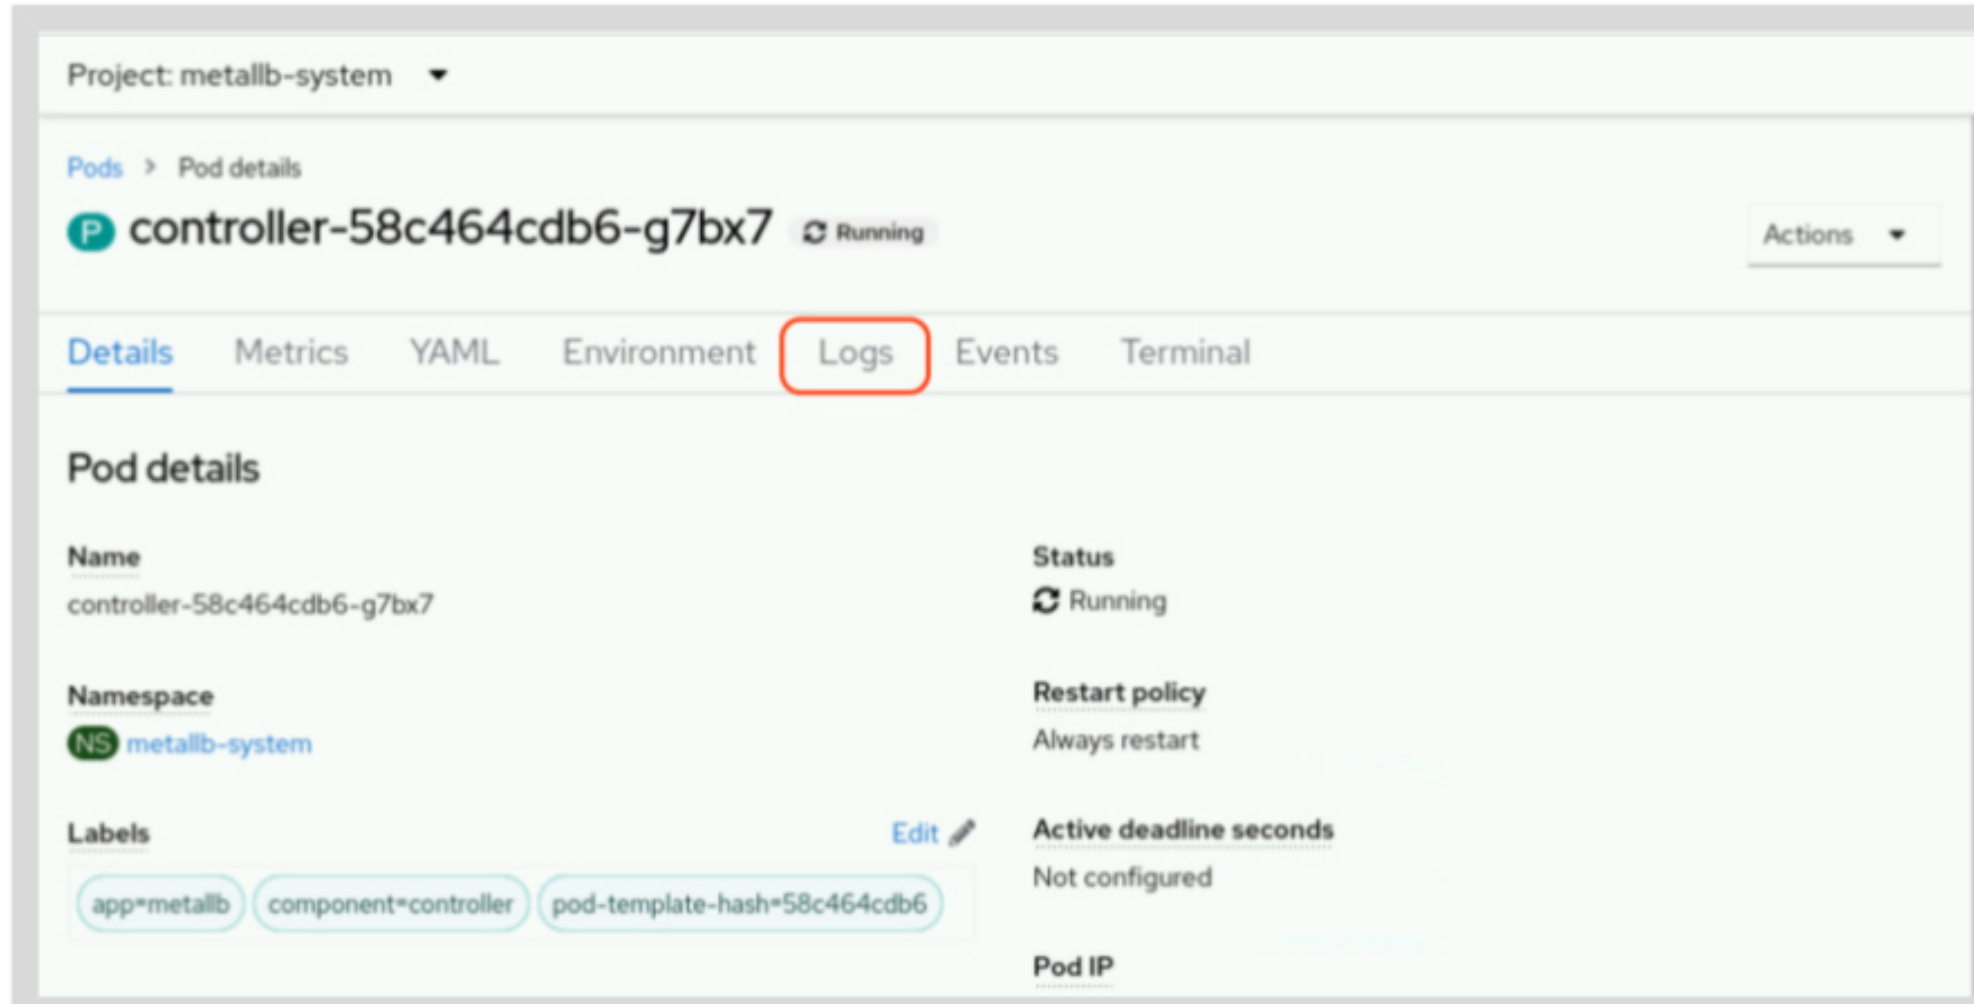

Access Container Logs using command

- Standard output (stdout) and standard error (stderr)
- Retrieve log using **oc logs** command, use **--tail=10** to limit lines of output

```
[user@host ~]$ oc logs postgresql-1-jw89j --tail=10
done
server stopped
Starting server...
2023-01-04 22:00:16.945 UTC [1] LOG:  starting PostgreSQL 12.11 on x86_64-redhat-
linux-gnu, compiled by gcc (GCC) 8.5.0 20210514 (Red Hat 8.5.0-10), 64-bit
2023-01-04 22:00:16.946 UTC [1] LOG:  listening on IPv4 address "0.0.0.0", port
5432
2023-01-04 22:00:16.946 UTC [1] LOG:  listening on IPv6 address ":::", port 5432
2023-01-04 22:00:16.953 UTC [1] LOG:  listening on Unix socket "/var/run/
postgresql/.s.PGSQL.5432"
2023-01-04 22:00:16.960 UTC [1] LOG:  listening on Unix socket "/"
tmp/.s.PGSQL.5432"
2023-01-04 22:00:16.968 UTC [1] LOG:  redirecting log output to logging collector
process
2023-01-04 22:00:16.968 UTC [1] HINT:  Future log output will appear in directory
"log".
```

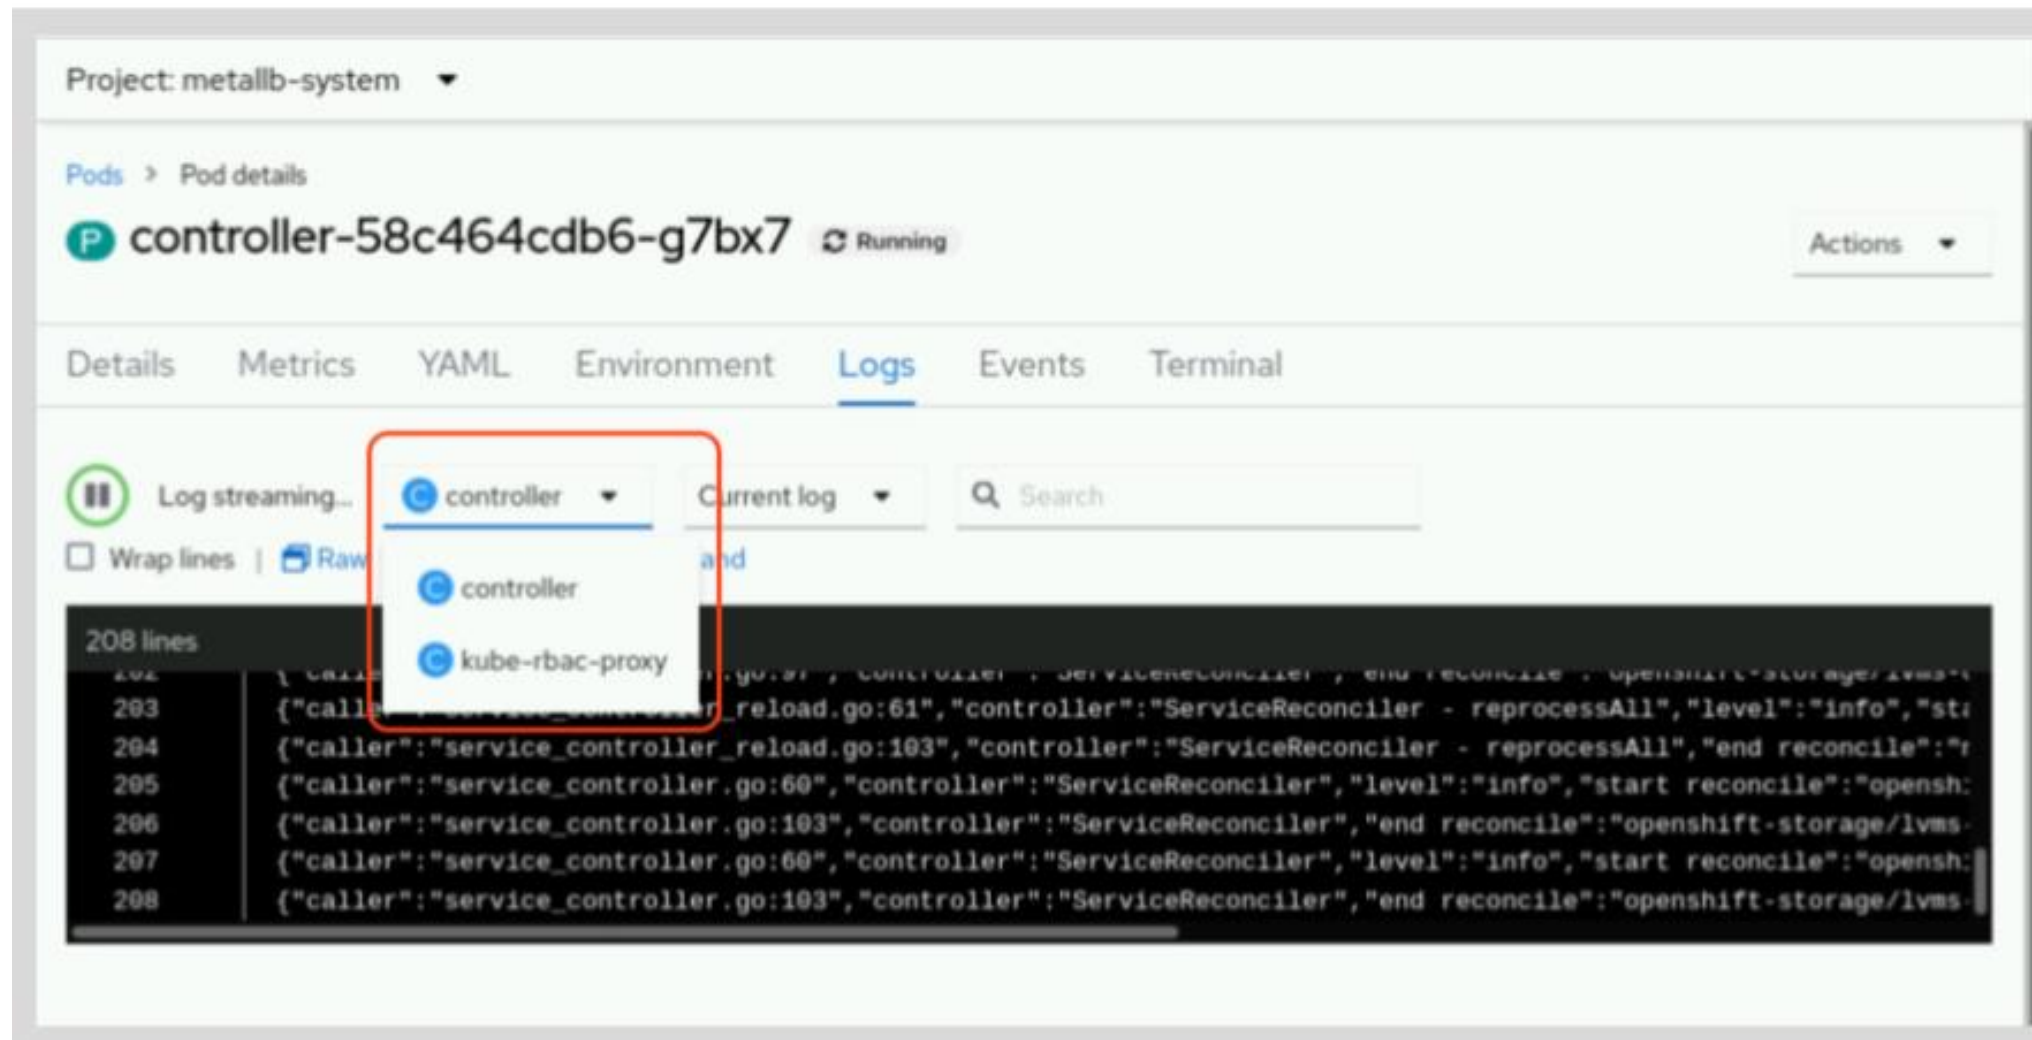
Access Container Logs using web console

- Access Pod then click on Logs tab



Access Container Logs using web console

- If there are multiple container, choose from drop down list
- Perform search, pause the stream



Deleting Resources

- Delete by resource type

```
[user@host ~]$ oc delete pod php-app
```

- Delete by label

```
[user@host ~]$ kubectl delete pod -l app=my-app  
pod "php-app" deleted  
pod "mysql-db" deleted
```

- Delete by resource type and a given json file

```
[user@host ~]$ oc delete pod -f ~/php-app.json  
pod "php-app" deleted
```

```
[user@host ~]$ cat ~/php-app.json | kubectl delete -f -  
pod "php-app" deleted
```

Deleting Resources

- Specify graceful period

```
[user@host ~]$ oc delete pod php-app --grace-period=10
```

- To shutdown pod immediately

```
[user@host ~]$ oc delete pod php-app --grace-period=10
```

- or use --now option

```
[user@host ~]$ oc delete pod php-app --now
```

- Forcibly delete pod

```
[user@host ~]$ kubectl delete pod php-app --force
```

- Delete all pods in project/namespace

```
[user@host ~]$ kubectl delete pods --all  
pod "php-app" deleted  
pod "mysql-db" deleted
```

Deleting Project

- Will remove all resources within project

```
[user@host ~]$ oc delete project my-app  
project.project.openshift.io "my-app" deleted
```

- `--timeout=#s` : specify length of time to wait before giving up on delete.
- `--now` : resources are signaled for immediate shutdown
- `--force` : no graceful period. Immediate deletion of some resources and may result in inconsistency or data lose

The CRI-O Container Engine

- is required to run containers
- Used in Worker and Control plane nodes in OpenShift
- Designed and optimized to run containers in Kubernetes cluster
- Meets Kubernetes CRI standards
 - can integrate resources with each other
 - such as networking and storage plug-ins

CRI-O management command

crictl pods

Lists all pods on a node.

crictl image

Lists all images on a node.

crictl inspect

Retrieve the status of one or more containers.

crictl exec

Run a command in a running container.

crictl logs

Retrieve the logs of a container.

crictl ps

List running containers on a node.

CRI-O management command - setup

- Find out which node did the pod runs

```
[user@host ~]$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
postgresql-1-8lzf2	1/1	Running	0	20m	10.8.0.64	master01
postgresql-1-deploy	0/1	Completed	0	21m	10.8.0.63	master01

- Create debug pod to access the node

```
[user@host ~]$ oc debug node/worker01
Starting pod/worker01-debug ...
To use host binaries, run chroot /host
Pod IP: 192.168.50.10
If you don't see a command prompt, try pressing enter.
sh-4.4# chroot /host
```

CRI-O management command - diagnose

- Find the container ID

```
sh-4.4# crictl ps --name postgresql
```

CONTAINER	IMAGE	CREATED	STATE	NAME	ATTEMPT	POD ID	POD
27943ae4f3024	image...7104	5...ago	Running	postgresql	0	5768...f015	

```
postgresql-1...

sh-4.4# crictl ps --name postgresql -o json | jq .containers[0].id
"2794...29a4"
```

- Retrieve PID of running container

```
sh-4.4# crictl inspect -o json 27943ae4f3024 | jq .info.pid
43453
sh-4.4# crictl inspect 27943ae4f3024 | grep pid
    "pid": 43453,
...output omitted...
```

CRI-O management command - diagnose

- List the system namespaces of the container

```
sh-4.4# lsns -p 43453
```

	NS	TYPE	NPROCS	PID	USER	COMMAND
4026531835	cgroup		530	1	root	/usr/lib/systemd/systemd --switched-root
		--system				--deserialize 17
4026531837	user		530	1	root	/usr/lib/systemd/systemd --switched-root
		--system				--deserialize 17
4026537853	uts		8	43453	1000690000	postgres
4026537854	ipc		8	43453	1000690000	postgres
4026537856	net		8	43453	1000690000	postgres
4026538013	mnt		8	43453	1000690000	postgres
4026538014	pid		8	43453	1000690000	postgres

Troubleshooting a latency-sensitive application in container - use of nsenter

- Use nsenter to execute a command within specified namespace of running container

```
sh-4.4# nsenter -t 43453 -p -r ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
1000690+	1	0	0	18:49	?	00:00:00	postgres
1000690+	58	1	0	18:49	?	00:00:00	postgres: logger
1000690+	60	1	0	18:49	?	00:00:00	postgres: checkpointer
1000690+	61	1	0	18:49	?	00:00:00	postgres: background writer
1000690+	62	1	0	18:49	?	00:00:00	postgres: walwriter
1000690+	63	1	0	18:49	?	00:00:00	postgres: autovacuum launcher
1000690+	64	1	0	18:49	?	00:00:00	postgres: stats collector
1000690+	65	1	0	18:49	?	00:00:00	postgres: logical replication launcher
root	7414	0	0	20:14	?	00:00:00	ps -ef

Troubleshooting a latency-sensitive application in container - use of nsenter

- Use of -a option to execute a command in all the container's namespaces

```
sh-4.4# nsenter -t 43453 -a ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
1000690+	1	0	0	18:49	?	00:00:00	postgres
1000690+	58	1	0	18:49	?	00:00:00	postgres: logger
1000690+	60	1	0	18:49	?	00:00:00	postgres: checkpointer
1000690+	61	1	0	18:49	?	00:00:00	postgres: background writer
1000690+	62	1	0	18:49	?	00:00:00	postgres: walwriter
1000690+	63	1	0	18:49	?	00:00:00	postgres: autovacuum launcher
1000690+	64	1	0	18:49	?	00:00:00	postgres: stats collector
1000690+	65	1	0	18:49	?	00:00:00	postgres: logical replication launcher
root	10058	0	0	20:45	?	00:00:00	ps -ef

Guided Exercise: Create Linux Containers and Kubernetes Pods

You should be able to:

- Create a pod
- View the logs of a running container.
- Retrieve information inside a container,
- Identify the process ID (PID) and namespaces for a container.
- Identify the User ID (UID) and supplemental group ID (GID) ranges of a project.
- Compare the namespaces of containers in one pod versus in another pod.
- Inspect a pod with multiple containers, and identify the purpose of each container.

In this chapter, you learned:



The OpenShift web console provides a GUI for visualizing and managing OpenShift resources.



Some resources feature a specialized page that makes creating and editing resources more convenient than writing YAML by hand, such as the Edit Key/Value Secret editor, which automatically handles Base64 encoding and decoding.



You can install partner and community operators from the embedded OperatorHub page.



Cluster-wide metrics such as CPU, memory, and storage usage are displayed on the Dashboards page.



Project Details pages display metrics specific to the project, such as the top ten memory consumers by pod and the current resource quota usage.