

Lesson 4 : Manage Non-shared Storage with Stateful Sets

Deploy applications that scale without sharing storage.

Application Clustering

- Clustering applications requires persistent storage
 - MySQL
 - Cassandra
- Multiple applications requires access to database system
- Need strong or eventual consistency read/write
- Common setup: Shared storage via NAS

Network Attached Storage (NAS)

- File-based storage
- Good for data backup, archiving
- Has reliability, fault-tolerant, security, standard protocol, lock handler and caching
- No need for proprietary hardware / software
- Common protocol: NFS and CIFS (SMB)
- Common case uses:
 - Web content
 - File share services
 - FTP storage
 - Backup archives
 - Video streaming, audio streaming
- Main disadvantage: not fit for structurable content

Storage Area Network (SAN)

- Block-level storage
- Mostly required proprietary hardware or connections
- Protocols:
 - FC, FCoE, iSCSI, NVMeoFC, NVMeoTCP
- Present raw device / LUN to hosts
- LUNs sharing; application handle locks and release
- Data protection: RAID configuration
- Common use:
 - SQL Databases (single node access), VM (multimode access)
 - High-performance data access, Server-side processing applications
 - Clustering applications and many more

Stateless vs Statefull

Stateless

- Request sent to server and response relayed back without storing any information
- Or Data loss upon user exit/logouts
- Less complex setup
- Typically slower than Stateful
- Every transaction recreate, repeat, recalculate

Statefull

- Store data persistently
- Data can be used/tracked by servers, clients, other applications
- Simplifies data recovery (point in time)
- Requires design to complex setup
- Faster processing in subsequent transactions

Stateful Sets

- Stateful application simplifies recovery from failures
- **Consistent identities**: Single stable DNS, hostname, storage
- Stateful set **guarantees given network identity maps to same storage identity**
 - **Consistent permission** to volumes, network resources
 - Facilitates **graceful** scaling operations and rolling updates
 - Pods added and removed in predictable order
- Best options for applications (databases) require **consistent identities** and **non-shared persistent** storage

Differences

- Deployments provides **stateless** representation of set of pods
- StatefulSets provides **stateful** representation of set of pods
- ReplicaSets help **manage traffic** by scaling and load-balance multi-pods application

Example application

Database requirements:

mysql-0 – First pod, primary role (read-write) → sit on higher end node

mysql-1 – Read-only replica → sits on lower end node

mysql-2 – Read-only replica → sits on lower end node

- Applications / user request connect to primary role for read-write access
- Event of failure, next replica is promoted to primary role, continue access by application
- However when first replica back online, failback can be automated
- If scaleDown occurs, only read-only replicas removed first

Create StatefulSets using YAML manifest file

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: dbserver ❶
spec:
  selector:
    matchLabels:
      app: database ❷
  replicas: 3 ❸
  template:
    metadata:
      labels:
        app: database ❹
    spec:
      containers:
        - env: ❺
          - name: MYSQL_USER
            valueFrom:
              secretKeyRef:
                key: user
                name: sakila-cred
          image: registry.ocp4.example.com:8443/redhattraining/mysql-app:v1 ❻
          name: database ❼
          ports: ❽
          - containerPort: 3306
            name: database
          volumeMounts: ❾
          - mountPath: /var/lib/mysql
            name: data
      terminationGracePeriodSeconds: 10
```

1. Name of the stateful set.

2,4 Application labels.

3 Number of replicas.

5 Environment variables (can use secret object).

6 Image source.

7,8 Container name and ports.

9 Mount path information for the persistent volumes for each replica. Each persistent volume has the same configuration.

Create StatefulSets using YAML manifest file

Create the stateful set by using the create command:

```
[user@host ~]$ oc create -f statefulset-dbserver.yml
```

Verify the creation of the stateful set named dbserver:

```
[user@host ~]$ kubectl get statefulset
```

NAME	READY	AGE
dbserver	3/3	6s

Verify the status of the instances:

```
[user@host ~]$ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
dbserver-0	1/1	Running	0	85s
dbserver-1	1/1	Running	0	82s
dbserver-2	1/1	Running	0	79s

Verify the status of the persistent volumes:

```
[user@host ~]$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS	MODES
data-dbserver-0	Bound	pvc-c28f61ee-...	1Gi	RWO	nfs-
data-dbserver-1	Bound	pvc-ddbe6af1-...	1Gi	RWO	nfs-
data-dbserver-2	Bound	pvc-8302924a-...	1Gi	RWO	nfs-

Notice that three PVCs were created. Confirm that persistent volumes are attached to each instance:

```
[user@host ~]$ oc describe pod dbserver-0
...output omitted...
Volumes:
  data:
    Type:      PersistentVolumeClaim (a reference
the same namespace)
    ClaimName: data-dbserver-0
...output omitted...
```

```
[user@host ~]$ oc describe pod dbserver-1
...output omitted...
Volumes:
  data:
    Type:      PersistentVolumeClaim (a ref
the same namespace)
    ClaimName: data-dbserver-1
...output omitted...
```

```
[user@host ~]$ oc describe pod dbserver-2
...output omitted...
Volumes:
  data:
    Type:      PersistentVolumeClaim (a re
the same namespace)
    ClaimName: data-dbserver-2
...output omitted...
```

StatefulSet lifecycles

You can update the number of replicas of the stateful set by using the `scale` command:

```
[user@host ~]$ oc scale statefulset/dbserver --replicas 1
```

NAME	READY	STATUS	RESTARTS	...
dbserver-0	1/1	Running	0	...

To delete the stateful set, use the `delete statefulset` command:

```
[user@host ~]$ kubectl delete statefulset dbserver
statefulset.apps "dbserver" deleted
```

Notice that the PVCs are not deleted after the execution of the `oc delete statefulset` command:

```
[user@host ~]$ oc get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS	MODES
data-dbserver-0	Bound	pvc-c28f61ee-...	1Gi	RWO	nfs-
data-dbserver-1	Bound	pvc-ddbe6af1-...	1Gi	RWO	nfs-
data-dbserver-2	Bound	pvc-8302924a-...	1Gi	RWO	nfs-

You can create a stateful set from the web console by clicking the **Workloads > StatefulSets** menu. Click **Create StatefulSet** and customize the YAML manifest.

Guided Exercise: Manage Non-shared Storage with Stateful Sets

You should be able to:

- Deploy a web server with persistent storage.
- Add data to the persistent storage.
- Scale the web server deployment and observe the data that is shared with the replicas.
- Create a database server with a stateful set by using a YAML manifest file.
- Verify that each instance from the stateful set has a persistent volume claim.

Lab: Manage Storage for Application Configuration and Data

You should be able to:

- Deploy a database server.
- Deploy a web application.
- Create a secret that contains the database server credentials.
- Create a configuration map that contains an SQL file.
- Add and remove a volume on the database server and the web application.
- Expose the database server and the web application.
- Scale up the web application.
- Mount the configuration map as a volume.